

# Concurrent Programming

Annalise Tarhan

October 16, 2020

## 12 Homework Problems

**12.16** Write a version of `hello.c` that creates and reaps `n` joinable peer threads, where `n` is a command-line argument.

```
int main(int argc, char *argv[])
{
    if (argc != 2) {
        fprintf(stderr, "usage: %s \n", argv[0]);
        exit(0);
    }

    int num_threads = atoi(argv[1]);
    pthread_t tids[num_threads];

    for (int i = 0; i < num_threads; i++) {
        pthread_t tid;
        pthread_create(&tid, NULL, thread, NULL);
        tids[i] = tid;
    }

    for (int i = 0; i < num_threads; i++) {
        pthread_join(tids[i], NULL);
    }

    exit(0);
}
```

## 12.17

- 12.17.1** The program below has a bug. The thread is supposed to sleep for 1 second and then print a string. However, when we run it on our system, nothing prints. Why?

When one thread in a process exits, all the other threads in that process are killed. The sleep call guarantees that the first thread will exit before the string is printed.

- 12.17.2** You can fix this bug by replacing the exit function with one of two different pthreads function calls. Which ones?

pthread\_join and pthread\_detach

- 12.18** Using the progress graph, classify the following trajectories as either safe or unsafe.

- 12.18.1**  $H_2, L_2, U_2, H_1, L_1, S_2, U_1, S_1, T_1, T_2$

unsafe

- 12.18.2**  $H_2, H_1, L_1, U_1, S_1, L_2, T_1, U_2, S_2, T_2$

safe

- 12.18.3**  $H_1, L_1, H_2, L_2, U_2, S_2, U_1, S_1, T_1, T_2$

unsafe

**12.19** Derive a solution to the first readers-writers problem that gives stronger priority to readers, where a writer leaving its critical section will always restart a waiting reader if one exists.

Use an additional semaphore for writers. In addition to the mutex semaphore that protects readcnt and the w semaphore that writers the first reader in lock, add an extra W semaphore that only writers lock and unlock. Writers must lock W before they can lock w, and they unlock w before they unlock W. Since readers don't have to wait for W to be unlocked, they can lock w before another writer is able to.

```
sem_t mutex, w, W;

void writer(void)
{
    while (1) {
        P(&W);
        P(&w);
        ...
        V(&w);
        V(&W);
    }
}
```

**12.20 Derive a solution to the readers-writers problem where there are at most N readers that gives equal priority to readers and writers, in the sense that pending readers and writers have an equal chance of being granted access to the resource.**

Credit due elsewhere. With the limit on the number of readers, access becomes a scarce resource, even for readers. They take one slot at a time, but this solution allows a waiting writer to claim the slots as well. Once the writer has acquired all N slots, it is free to write, since all of the readers are locked out. When it is finished, it releases all N slots. Writers don't compete, since there is a mutex guarding the slot-claiming routine.

```
int N;
sem_t mutex; /* Initially = 1 */
sem_t slots; /* Initially = N */

void reader(void)
{
    P(&slots);
    /* Critical section */
    V(&slots);
}

void writer(void)
{
    P(&mutex);
    for (int i = 0; i < N; i++) {
        P(&slots);
    }
    /* Critical section */
    for (int i = 0; i < N; i++) {
        V(&slots);
    }
    V(&mutex);
}
```

**12.21    Derive a solution to the second readers-writers problem, which favors writers over readers.**

This solution is inefficient, since it doesn't allow more than one reader at a time. It sets up two semaphores for readers and only one for writers. Only one of any kind is allowed access to the resource at a time, and since the readers have to pass an extra hurdle, it favors writers.

```
void reader (void)
{
    P(&reader_gate);
    P(&mutex);
    /* Critical section */
    V(&mutex);
    V(&reader_gate);
}

void writer (void)
{
    P(&mutex);
    /* Critical section */
    V(&mutex);
}
```

## 12.22 Modify the server so that it echoes at most one text line per iteration of the main server loop.

Instead of connecting and reading from `connfd` directly, this version connects and adds `connfd` to the read set. On the subsequent iteration, it will start reading. It also calls `echo_line` instead of `echo`, which returns 1 if it reached EOF and 0 otherwise. When `echo_line` returns 1, it removes `connfd` from the read set and closes the connection. Otherwise, it carries on to the next iteration and reads another line.

```
int has_open_connection = 0;

while (1) {
    ready_set = read_set;
    select(listenfd+2, &ready_set, NULL, NULL, NULL);
    if (FD_ISSET(STDIN_FILENO, &ready_set))
        command();
    if (has_open_connection) {
        if (echo_line(connfd)) {
            close(connfd);
            FD_CLEAR(connfd, &read_set);
            has_open_connection = 0;
        }
    }
    if (!has_open_connection
        && FD_ISSET(listenfd, &ready_set)) {
        clientlen =
            sizeof(struct sockaddr_storage);
        connfd = accept(listenfd,
            (SA *)&clientaddr, &clientlen);
        FD_SET(connfd, &read_set);
        has_open_connection = 1;
    }
}
```

## 12.23 The event-driven echo server is flawed because a malicious client can deny service to other clients by sending a partial text line. Write an improved version of the server that can handle these partial text lines without blocking.

I suppose it would be cheating to throw each readline in its own thread?

**12.24 The functions in the RIO I/O package are thread-safe. Are they reentrant as well?**

They are reentrant, but only implicitly since some of the arguments are pointers to what could be shared data.

**12.25 In the prethreaded concurrent echo server, each thread calls the echo\_cnt function. Is echo\_cnt thread-safe? Is it reentrant? Why or why not?**

The function is thread-safe but not reentrant. It isn't reentrant because it references static global variables. It uses a semaphore to protect shared variables, though, so it is thread-safe.

**12.26 Use the lock-and-copy technique to implement a thread-safe non-reentrant version of gethostbyname called gethostbyname\_ts.**

```
struct hostent *gethostbyname_ts(const char *name)
{
    struct hostent *sharedp;
    struct hostent *copied_hostent = NULL;

    P(&mutex);
    sharedp = gethostbyname(name);
    copied_hostent->h_name = sharedp->h_name;
    copied_hostent->h_aliases = sharedp->h_aliases;
    copied_hostent->h_addrtype = sharedp->h_addrtype;
    copied_hostent->h_length = sharedp->h_length;
    copied_hostent->h_addr_list = sharedp->h_addr_list;
    V(&mutex);
    return copied_hostent;
}
```

**12.27 Why does the following approach for reading and writing sockets, which opens two standard I/O streams on the same open connected socket descriptor, create a deadly race condition in a concurrent server based on threads?**

Each fclose operation attempts to close the same underlying socket descriptor. In a sequential program, this isn't a problem, but in a concurrent one that socket descriptor could have already been assigned to a different, newly opened socket and the second fclose operation could close that socket instead.

**12.28 Does swapping the order of the two V operations have any effect on whether or not the program deadlocks?**

No it doesn't. It restricts the possible trajectories through the state space, but none of the orderings lead to deadlock. Deadlock would only occur if the threads' P operations were in different orders.

**12.29 Can the following program deadlock? Why or why not?**

No it can't. By the mutex lock ordering rule, the program is deadlock-free since each thread acquires mutexes in order and releases them in reverse order.

**12.30 Consider the following program that deadlocks.**

**12.30.1 For each thread, list the pairs of mutexes that it holds simultaneously.**

**Thread 1:** (a,b) (a, c)

**Thread 2:** (c, b)

**Thread 3:** (b, a)

**12.30.2 If  $a < b < c$ , which threads violate the mutex lock ordering rule?**

Threads 2 and 3. Thread 2 acquires c before b and thread 3 acquires b before a. Technically, thread 1 releases b before acquiring c and threads 2 and 3 acquire c before anything else, but those infractions are less significant.

**12.30.3 For these threads, show a new lock ordering that guarantees freedom from deadlock.**

Given the order  $a < b < c$ , the perfect lock ordering is P(a); P(B); P(C); V(C); V(B); V(A).



**12.31** Implement a version of the standard I/O `fgets` function, called `tfgets`, that times out and returns `NULL` if it does not receive an input line on standard input within 5 seconds. It should use processes, signals, and nonlocal jumps.

```
jmp_buf buf;

void sigchld_handler(int sig)
{
    longjmp(buf, 1);
}

char *tfgets(char *s, int size, FILE *stream) {
    if (fork() == 0) {
        sleep(5);
        exit(0);
    }
    signal(SIGCHLD, sigchld_handler);
    if (setjmp(buf)) {
        return NULL;
    } else {
        return fgets(s, size, stream);
    }
}
```

**12.32** Implement another version of the `tfgets` function that uses the `select` function.

```
char *tfgets(char *s, int size, FILE *stream)
{
    struct timeval tv;
    tv.tv_sec = 5;
    tv.tv_usec = 0;

    fd_set read_set;
    FD_ZERO(&read_set);
    FD_SET(STDIN_FILENO, &read_set);

    if (select(1, &read_set, NULL, NULL, &tv)) {
        return fgets(s, size, stream);
    } else {
        return NULL;
    }
}
```

### 12.33 Implement a threaded version of the tfgets function.

```
int complete = 0;
char *result = NULL;

struct fgets_args {
    char *s;
    int size;
    FILE *stream;
};

void *thread1(void *vargp) {
    sleep(5);
    complete = 1;
    return NULL;
}

void *thread2(void *vargp) {
    struct fgets_args *args =
        (struct fgets_args *) vargp;
    result = fgets(args->s, args->size, args->stream);
    complete = 1;
    return NULL;
}

char *tfgets(char *s, int size, FILE *stream)
{
    struct fgets_args args;
    args.s = s;
    args.size = size;
    args.stream = stream;

    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, thread1, NULL);
    pthread_create(&tid2, NULL, thread2, &args);
    while (!complete) {}
    return result;
}
```

### 12.34 Write a parallel threaded version of an $N \times M$ matrix multiplication kernel. Compare the performance to the sequential case.

This version creates a new thread for each entry in the final  $N \times N$  matrix, which is extremely inefficient. For a  $10 \times 10$  matrix, the concurrent version ran almost a thousand times slower. A (much) better approach would have been to divide the matrix into around four parts and split the work between as many threads. Actually, this is probably a great example of code I'll be absolutely horrified to look at in a few years.

```
struct coordinates {
    int i;  int j;
};
void thread_mul()
{
    int i, j;
    pthread_t tids[10][10];
    struct coordinates crds[10][10];
    for (i = 0; i < 10; i++) {
        for (j = 0; j < 10; j++) {
            struct coordinates coords;
            coords.i = i;
            coords.j = j;
            crds[i][j] = coords;
            pthread_create(
                &tids[i][j], NULL,
                thread, &crds[i][j]);
        }
    }
    for (i = 0; i < 10; i++) {
        for (j = 0; j < 10; j++) {
            pthread_join(tids[i][j], NULL);
        }
    }
}
void *thread(void *vargp) {
    int i = ((struct coordinates *) vargp) -> i;
    int j = ((struct coordinates *) vargp) -> j;
    int sum = 0;
    for (int k = 0; k < 8; k++) {
        sum += A[i][k]*B[k][j];
    }
    C[i][j] = sum;
    return NULL;
}
```

**12.35 Implement a concurrent version of the TINY Web server based on processes. It should create a new child process for each new connection request.**

```
while (1) {
    clientlen = sizeof(clientaddr);
    connfd = Accept(
        listenfd, (SA *)&clientaddr, &clientlen);
    if (Fork() == 0) {
        Close(listenfd);
        Getnameinfo(
            (SA *)&clientaddr, clientlen,
            hostname, MAXLINE, port,
            MAXLINE, 0);
        doit(connfd);
        Close(connfd);
        exit(0);
    }
    Close(connfd);
}
```

**12.36 Implement a concurrent version of the TINY Web server based on I/O multiplexing.**

```
fd_set read_set, ready_set;
FD_ZERO(&read_set);
FD_SET(listenfd, &read_set);

while (1) {
    ready_set = read_set;
    Select(listenfd+1, &ready_set, NULL, NULL, NULL);
    if (FD_ISSET(listenfd, &ready_set)) {
        clientlen = sizeof(clientaddr);
        connfd = Accept(
            listenfd, (SA *)&clientaddr,
            &clientlen);
        Getnameinfo(
            (SA *)&clientaddr, clientlen,
            hostname, MAXLINE, port,
            MAXLINE, 0);
        doit(connfd);
        Close(connfd);
    }
}
```

**12.37 Implement a concurrent version of the TINY Web server based on threads.**

```
int *connfdp;
pthread_t tid;

while (1) {
    clientlen = sizeof(clientaddr);
    Getnameinfo(
        (SA *) &clientaddr, clientlen, hostname,
        MAXLINE, port, MAXLINE, 0);
    connfdp = malloc(sizeof(int));
    *connfdp = Accept(
        listenfd, (SA *)&clientaddr, &clientlen);
    Pthread_create(&tid, NULL, thread, connfdp);
}

void *thread(void *vargp)
{
    int connfd = *((int *)vargp);
    pthread_detach(pthread_self());
    free(vargp);
    doit(connfd);
    Close(connfd);
    return NULL;
}
```